

```
import pandas as pd

df = pd.read_csv('/content/survey_lung_cancer.csv')

print("First 5 rows of the dataset:")
display(df.head())

print("\nColumn names and their data types:")
display(df.info())
```

First 5 rows of the dataset:

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL CONSUMING	COUGHING	SHORTNESS OF BREATH
0	M	69	1	2	2	1	1	2	1	2	2	2	2
1	M	74	2	1	1	1	2	2	2	1	1	1	2
2	F	59	1	1	1	2	1	2	1	2	1	2	2
3	M	63	2	2	2	1	1	1	1	1	2	1	1
4	F	63	1	2	1	1	1	1	1	2	1	2	2

```
Column names and their data types:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   GENDER                                309 non-null    object
1   AGE                                   309 non-null    int64
2   SMOKING                              309 non-null    int64
3   YELLOW_FINGERS                       309 non-null    int64
4   ANXIETY                              309 non-null    int64
5   PEER_PRESSURE                        309 non-null    int64
6   CHRONIC DISEASE                      309 non-null    int64
7   FATIGUE                              309 non-null    int64
8   ALLERGY                              309 non-null    int64
9   WHEEZING                             309 non-null    int64
10  ALCOHOL CONSUMING                    309 non-null    int64
11  COUGHING                             309 non-null    int64
12  SHORTNESS OF BREATH                  309 non-null    int64
13  SWALLOWING DIFFICULTY                309 non-null    int64
14  CHEST PAIN                           309 non-null    int64
15  LUNG_CANCER                          309 non-null    object
dtypes: int64(14), object(2)
memory usage: 38.8+ KB
None
```

```
from sklearn.model_selection import train_test_split

X = df.drop('LUNG_CANCER', axis=1)
y = df['LUNG_CANCER']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)
```

```
Shape of X_train: (247, 15)
Shape of X_test: (62, 15)
Shape of y_train: (247,)
Shape of y_test: (62,)
```

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Shape of X_train_scaled:", X_train_scaled.shape)
print("Shape of X_test_scaled:", X_test_scaled.shape)
```



```
Shape of X_train_scaled: (247, 15)
Shape of X_test_scaled: (62, 15)
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

model = LogisticRegression()
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1:.4f}")
```

```
Accuracy: 0.9677
Precision: 0.9833
Recall: 0.9833
F1-score: 0.9833
```

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(3, 2))
sns.countplot(data=df, x='LUNG_CANCER')
```

<Axes: xlabel='LUNG_CANCER', ylabel='count'>



encoding

x and y me break

train test split

standard scaler

model *train*

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df['LUNG_CANCER'] = label_encoder.fit_transform(df['LUNG_CANCER'])
df['GENDER'] = label_encoder.fit_transform(df['GENDER'])

print(df.head())
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
0	1	69	1	2	2	1	
1	1	74	2	1	1	1	
2	0	59	1	1	1	2	
3	1	63	2	2	2	1	
4	0	63	1	2	1	1	

	CHRONIC_DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL_CONSUMING	COUGHING	\
0		1	2	1	2	2	2
1		2	2	2	1	1	1
2		1	2	1	2	1	2
3		1	1	1	1	2	1
4		1	1	1	2	1	2

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
0	2	2	2	1
1	2	2	2	1
2	2	1	2	0
3	1	2	2	0
4	2	1	1	0

```
from sklearn.svm import SVC

model=SVC()
model.fit(X_train,y_train)
```

▼ SVC ⓘ ?

SVC()

```
model.score(X_test,y_test)*100,model.score(X_train,y_train)*100

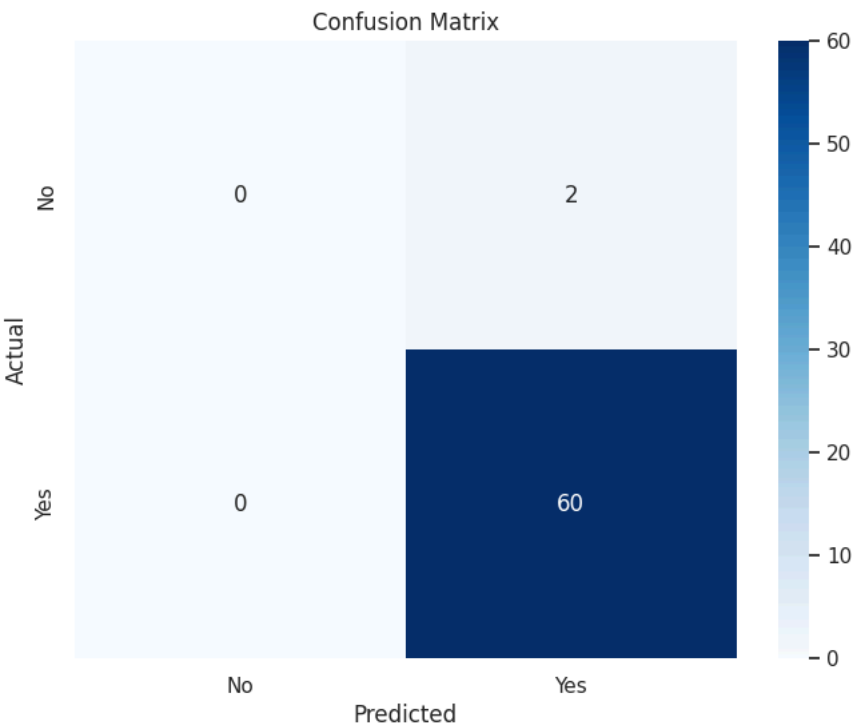
(96.7741935483871, 85.02024291497976)
```

```
y_pred=model.predict(X_test)
```

```
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
cm
```

```
array([[ 0,  2],
       [ 0, 60]])
```

```
# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	2

	1	0.97	1.00	0.98	60
accuracy				0.97	62
macro avg		0.48	0.50	0.49	62
weighted avg		0.94	0.97	0.95	62

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined a
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined a
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined a
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

# Get categorical columns
categorical_cols = df.select_dtypes(include=['object', 'category']).columns

```

```

# Create bar plots for categorical features
for col in categorical_cols:
    print(f"\nAnalyzing column: {col}")
    plt.figure(figsize=(8, 4))
    sns.countplot(data=df, x=col, order=df[col].value_counts().index)
    plt.title(f'Count of {col}')
    plt.xlabel(col)
    plt.ylabel('Count')
    plt.show()

```

```


sns.pairplot(df[['AGE', 'SMOKING', 'YELLOW_FINGERS', 'ANXIETY', 'LUNG_CANCER']], hue='LUNG_CANCER')
plt.suptitle('Pairwise relationships between selected features', y=1.02)
plt.show()

```

Pairwise relationships between selected features




```
plt.figure(figsize=(8, 6))
sns.histplot(data=df, x='AGE', kde=True)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```




```
sns.barplot(x='LUNG_CANCER', y='AGE', data=df)
```

<Axes: xlabel='LUNG_CANCER', ylabel='AGE'>




```
plt.figure(figsize=(8, 6))
sns.boxplot(data=df, x='LUNG_CANCER', y='AGE')
```

```
plt.title('Distribution of Age by Lung Cancer Outcome')  
plt.xlabel('Lung Cancer')  
plt.ylabel('Age')  
plt.show()
```



```
plt.figure(figsize=(12, 10))  
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")  
plt.title('Correlation Heatmap of Numerical Features')  
plt.show()
```



```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1:.4f}")
```

SWALLOWING DIFFICULTY
Accuracy: 0.9677
Precision: 0.9677
Recall: 1.0000
F1-score: 0.9836

CHEST PAIN

LUNG_CANCER

A continuation of the correlation heatmap, showing relationships between the seven numerical features and three categorical features: SWALLOWING DIFFICULTY, CHEST PAIN, and LUNG_CANCER. The diagonal elements are all 1.00. The heatmap uses a color scale from blue (negative correlation) to red (positive correlation). The values range from -0.16 to 0.49.

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL CONSUMING	COUGHING	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
SWALLOWING DIFFICULTY	-0.08	-0.00	0.03	0.35	0.49	0.37	0.08	-0.13	-0.06	0.07	-0.01	-0.16	-0.16	1.00	0.07	0.26
CHEST PAIN	0.36	-0.02	0.12	-0.10	-0.11	-0.09	-0.04	-0.01	0.24	0.15	0.33	0.08	0.02	0.07	1.00	0.19
LUNG_CANCER	0.07	0.09	0.06	0.18	0.14	0.19	0.11	0.15	0.33	0.25	0.29	0.25	0.06	0.26	0.19	1.00